



# The Last Man Standing:

The Only Practical, Lightweight and Hypervisor-Based  
Kernel Protector Struggling with the Real World Alone

Seunghun Han

[hanseunghun@nsr.re.kr](mailto:hanseunghun@nsr.re.kr)

# Who Am I?



- Senior security researcher at NSR (National Security Research Institute of South Korea)
- Speaker at
  - USENIX Security 2018
  - Black Hat Asia 2017, 2018
  - HITBSecConf 2016, 2017
  - KIMCHICON 2018
- Author of the book series titled “64-bit multi-core OS principles and structure, Vol.1&2”
- a.k.a kkamagui, @kkamagui1

# Who Am I?

## Offensive Researches

- IRON-HID: Create Your Own Bad USB Device
- I Don't Want to Sleep Tonight: Subverting Intel TXT with S3 Sleep
- A Bad Dream: Subverting Trusted Platform Module while You Are Sleeping

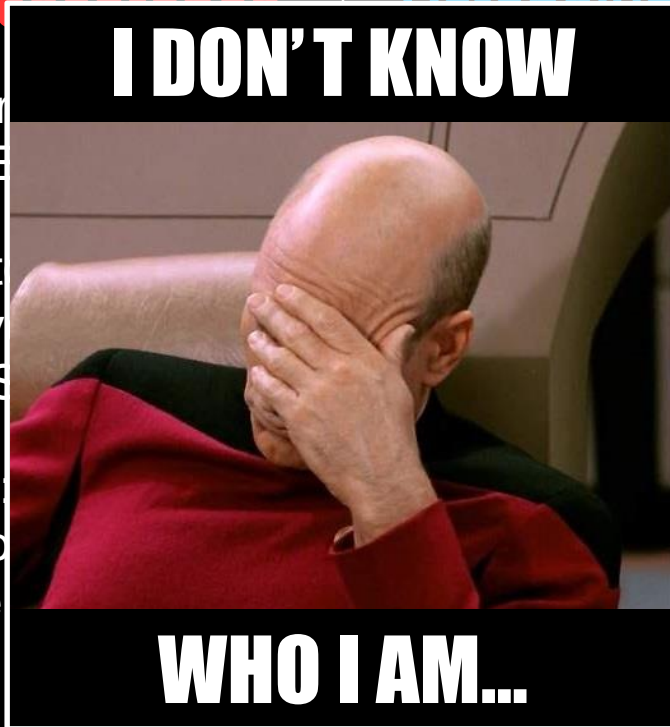
## Defensive Researches

- Myth and Truth about Hypervisor-Based Kernel Protector: The Reason Why You Need Shadow-Box
- Shadow-Box v2: The Practical and Omnipotent Sandbox for ARM
- Linux Kernel Vulnerability Patches

# Who Am I?

## Offensive Researches | Defensive Researches

- IRON-HID: Cracking the Code of Own Bad USB
- I Don't Want to Be a Hero Tonight: Subverting the TXT with S3
- A Bad Dream: The Trusted Platform Module while You Are



Truth about  
or-Based Kernel  
The Reason  
Need Shadow-

Box v2: The  
and Omnipotent  
for ARM

Kernel Vulnerability



**Background**

**Analysis of Privilege Escalation  
Exploits**

**Design and Implementation of  
Gatekeeper**  
(with Shadow-Box)

**Evaluation, Demo., Hunting the Beasts**  
(and Why I am Struggling with the Real World Beasts Alone)

**Limitations and Conclusion**

# **Background**

## **Analysis of Privilege Escalation Exploits**

## **Design and Implementation of Gatekeeper (with Shadow-Box)**

## **Evaluation, Demo., Hunting the Beasts (and Why I am Struggling with the Real World Beasts Alone)**

## **Limitations and Conclusion**



# Linux Kernel Is Everywhere!



# We Love Root!

Login **Linux**

root

Guest Session

Remote Login

Android

```
-- Install /zcard...
Finding update...
Opening update...
Verifying update package...
Installing update...
Rooting with SuperSU 0.95
Installing temporary busybox
Mounting system...
Removing old superuser
Installing Superuser...
Setting Permission...
Symlinking...
Unmounting system...
Deleting temporary busybox
Root complete!
SuperSU 0.95
```

Smart TV



IP Camera



T-Shirts!!





# **Exploit Vulnerabilities for Root!**

**Write Once, Run Anywhere!**

**- Java**

**Write Once, Run Anywhere and Pray!**

**- You and I**

# But, Our Prayers Went to **/dev/null**

- Kernel Address Space Layout Randomization (KASLR)
  - Makes exploits difficult to run Return-Oriented Programming (ROP) gadgets!
- Supervisor Mode Access Prevention (SMAP)
  - Makes **OUR** exploits hard to access user-level data!
- Supervisor Mode Execution Protection (SMEP)
  - Makes **OUR PRECIOUS** exploits super hard to execute user-level code!

# But, Our Prayers Went to **/dev/null**

- Kernel Address Space Layout Randomization (KASLR)
  - Makes exploits difficult to run Return-Oriented Programing (ROP) gadgets!
- Supervisor Mode Access Protection
  - Makes **OUR** exploits hard to execute on user-level data!
- Supervisor Mode Execution Protection (SMEP)
  - Makes **OUR PRECIOUS** exploits hard to execute user-level code!



# But, We Finally Got Root!

- KASLR
  - Is neutralized by **information leak vulnerabilities!**
- SMAP
  - Is defeated by **turning off SMAP bit of the CR4 register!**
- SMEP
  - Is also broken by **turning off SMEP bit of the CR4 register!**

# But, We Finally Got Root!

- KASLR
  - Is neutralized by **information leak vulnerabilities!**
- SMAP
  - Is defeated by **turning off the CR4 register!**
- SMEP
  - Is also broken by **turning off the CR4 register!**





**Background**

# **Analysis of Privilege Escalation Exploits**

**Design and Implementation of  
Gatekeeper**  
(with Shadow-Box)

**Evaluation, Demo., Hunting the Beasts**  
(and Why I am Struggling with the Real World Beasts Alone)

**Limitations and Conclusion**



# Disclaimer (1)

- ***This presentation contains harmful information*** for offensive security researchers
  - Techniques for detecting abnormal privilege escalations are included!
- ***Evening party might be canceled*** if you listen to this presentation until the end
  - Because you could be busy to subvert my detection techniques or upgrade your exploits!

# Disclaimer (2)

- If ***you want to go to the party after the presentation***
  - Check the GitHub repository and see it!



[https://github.com/kkamagui/  
shadow-box-for-x86](https://github.com/kkamagui/shadow-box-for-x86)

# Analysis of (Our) Privilege Escalation Exploits (1)

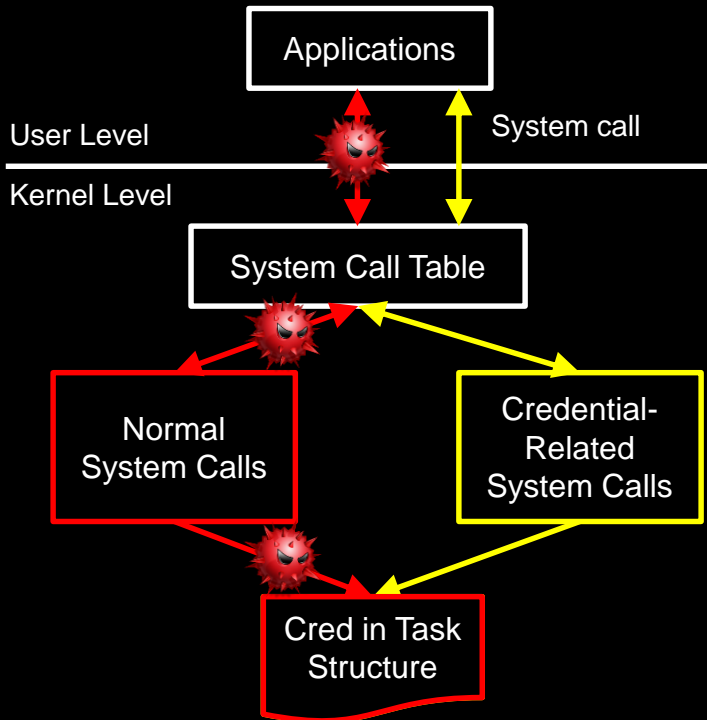
- (1) If vulnerabilities can read/write the kernel memory,
  - We overwrite **UID** in cred field of task structure with **ZERO**
  - ex) CVE-2017-16995: eBPF vulnerability
- (2) If vulnerabilities can hijack the control flow,
  - We make a **ROP chain** to execute below **“commit\_creds( prepare\_kernel\_cred(0))”!**
  - ex) CVE-2017-1000112: UDP fragmentation offload (UFO) vulnerability

# Analysis of (Our) Privilege Escalation Exploits (2)

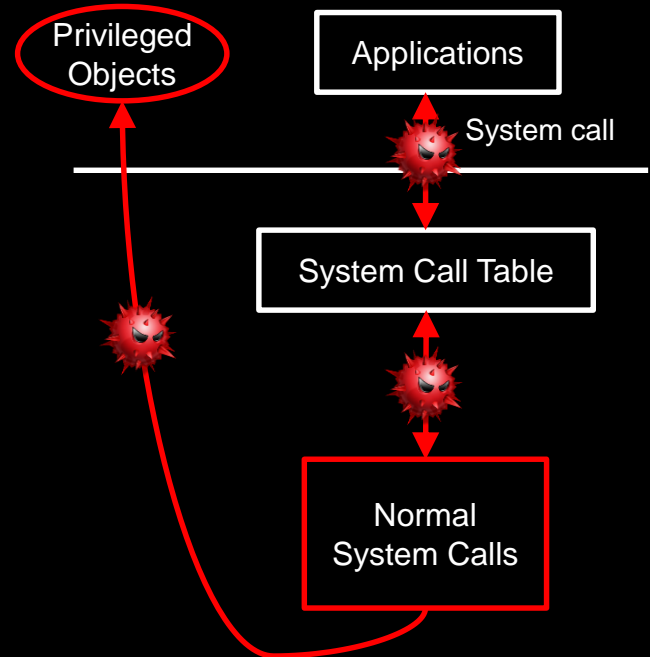
- (3) If vulnerabilities can read/write the privileged objects in user level,
  - We overwrite **/etc/crontab** or **/usr/bin/passwd** to get a root shell
  - ex) CVE-2016-4557: File descriptor Use-After-Free (UAF) vulnerability
  - ex) CVE-2016-5195: Copy-On-Write (COW) vulnerability a.k.a Dirty COW

# Call Flows of (Our) Privilege Escalation Exploits

## Call Flow of Case (1) and (2)



## Call Flow of Case (3)



# Detection and Prevention Techniques

- Case (1) and (2) change a cred of a task structure directly during normal system calls
  - Normal system calls don't change the credential
  - If someone **monitors system calls** and **changes o creds**, he can detect and prevent abnormal privilege escalations!
- Case (3) does not change a cred directly
  - It is hard to be detected and prevented!
  - If someone **monitors all processes created with UID = 0**, he can trace and detect it (not perfect, but possible)



# Detection and Prevention Techniques

- Case (1) and (2) change a cred of a task structure directly during normal system calls

**I made a tool and named**  
**“Gatekeeper”**

privilege escalations!

- Case (3) does not change a cred directly
  - It is hard to be detected and prevented!
  - If someone **monitors all processes created with UID = 0**, he can trace and detect it (not perfect, but possible)



**Background**

**Analysis of Privilege Escalation  
Exploits**

**Design and Implementation of  
Gatekeeper**  
(with Shadow-Box)

**Evaluation, Demo., Hunting the Beasts**  
(and Why I am Struggling with the Real World Beasts Alone)

**Limitations and Conclusion**

Simple Design

≠

Simple Implementation

Let's find something  
to **save** us

```
[ 10.161008] Bluetooth: RFCOMM ver 1.11
[ 11.796833] psmouse serio2: trackpoint: IBM TrackPoint firmware: 0x0e, buttons: 3/3
[ 11.993577] input: TPPS/2 IBM TrackPoint as /devices/platform/i8042/serio1/serio2/input/input9
[ 486.787670] shadow box: module verification failed: signature and/or required key missing - tainting kernel
[ 486.810413]
[ 486.810419]
[ 486.810421]
[ 486.810424]
[ 486.810426]
[ 486.810428]
[ 486.810430]
[ 486.810431]
[ 486.810433]
[ 486.810434]
```

# SHADOW-BOX

Lightweight Hypervisor-Based Kernel Protector

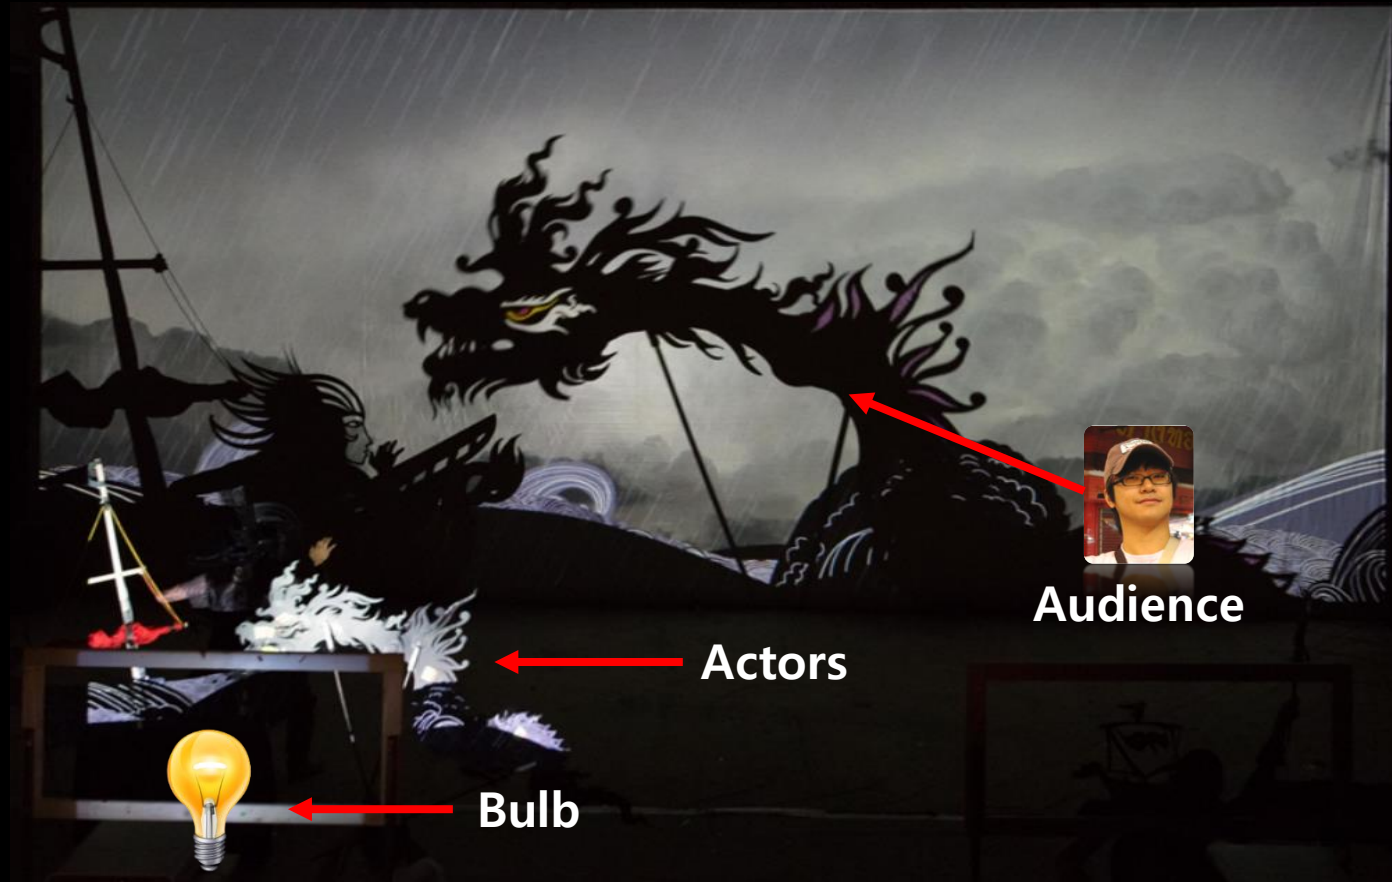
```
[ 486.826181] Shadow-box: CPU Count 4
[ 486.826185] Shadow-box: Booting CPU ID 0
[ 486.878993] Shadow-box: Protect Kernel Code Area
[ 486.904495] Shadow-box:      [*] Complete
[ 486.904497] Shadow-box: Protect Module Code Area
[ 486.910824] Shadow-box:      [*] Complete
[ 486.911132] Shadow-box: Framework Preinitialize
[ 486.912296] Shadow-box:      [*] Complete
[ 486.947846] Shadow-box: Framework Initailize
[ 486.947881] Shadow-box:      [*] Task count 216
[ 486.947882] Shadow-box:      [*] Module count 88
[ 486.947883] Shadow-box:      [*] Complete
[ 488.039134] Shadow-box: Lock IOMMU
[ 488.039732] Shadow-box:      [*] Lock IOMMU complete
[ 488.063655] Shadow-box: Execution Complete
[ 488.063658] Shadow-box: ErrorCode: 0
```

```
user@Shadow-Box:~/Demo$ █
```

# What Is Shadow-Box?

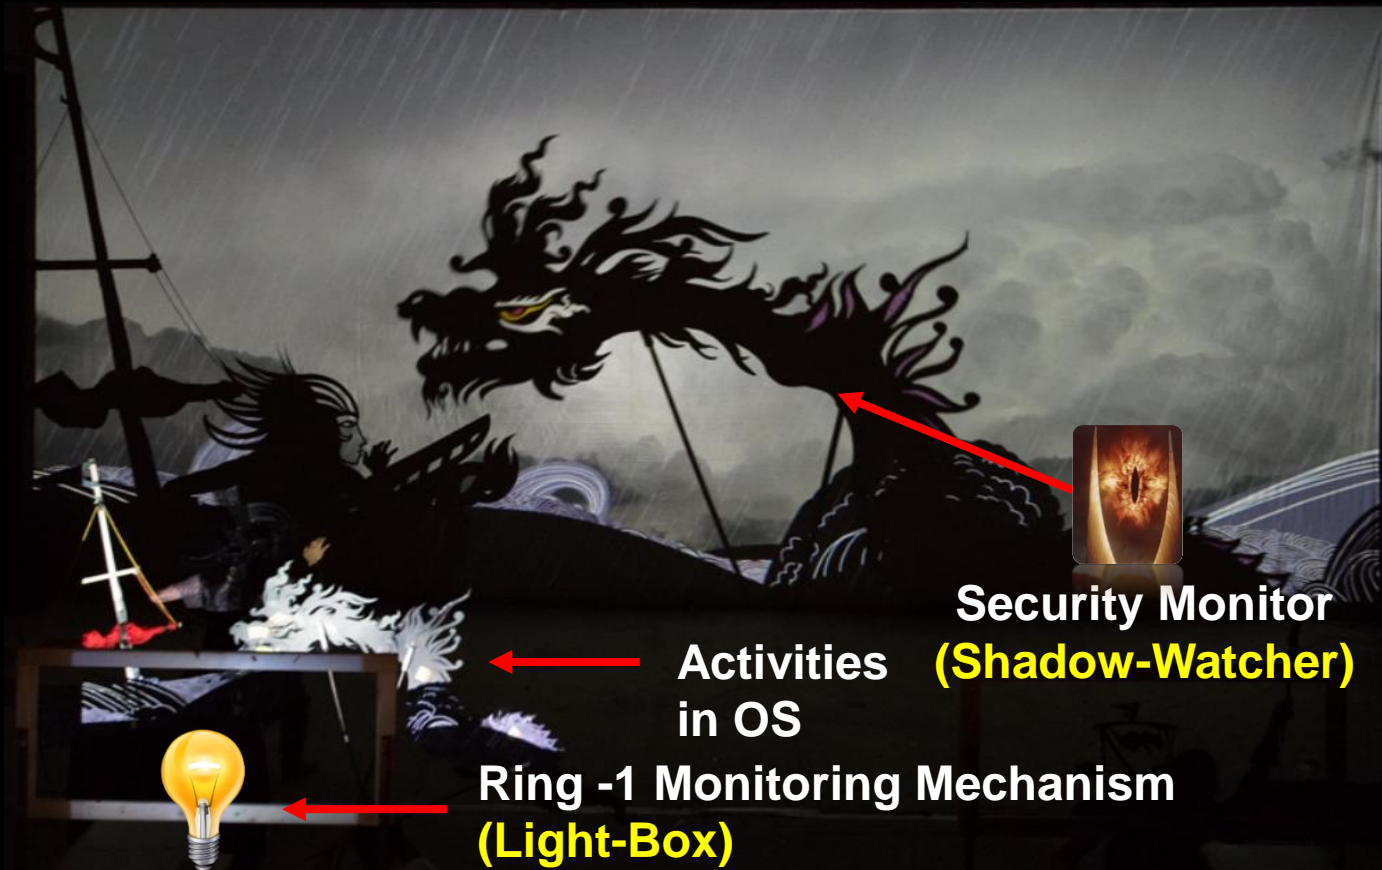
- Lightweight hypervisor-based kernel protector
  - Focuses on **ROOTKIT** detection and protection
  - has a small memory footprint
    - It shares kernel area and no multiple operating systems
- Practical and Portable
  - Out-of-box approach
    - No modification of the kernel code and data
  - Dynamic injection (Loadable Kernel Module, LKM)

# Security Architecture in Shadow Play





# Security Architecture in Shadow Play



# Security Architecture in Shadow Play

I named this architecture  
“**Shadow-box**”



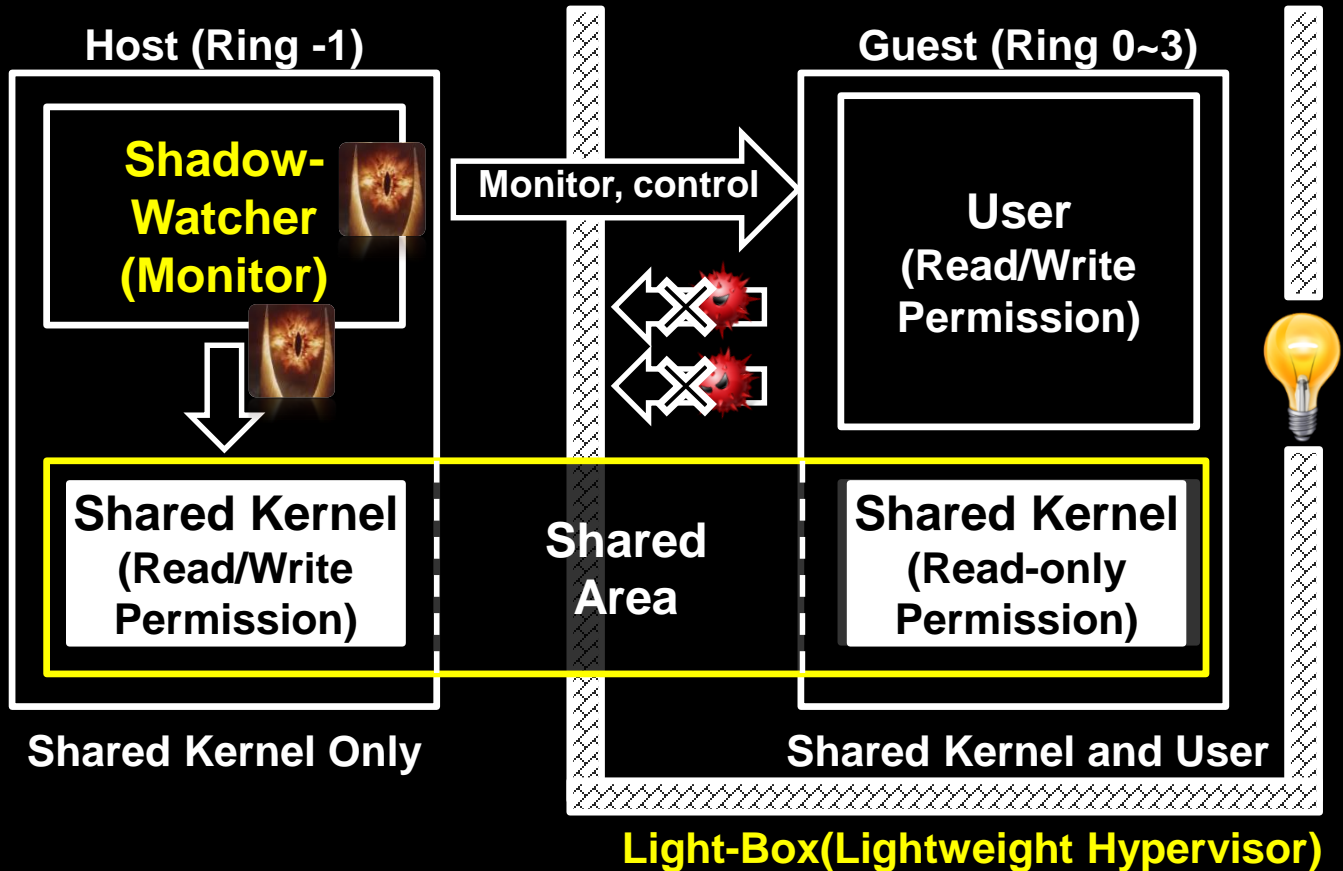
Security Monitor  
(**Shadow-Watcher**)

Activities  
in OS

Ring -1 Monitoring Mechanism  
(**Light-Box**)



# Architecture of Shadow-Box



# Architecture of Light-Box

- Light-box, lightweight hypervisor,
  - Isolates worlds with memory protection technique in Intel Virtualization Technology (VT-x and VT-d)
  - **Shares the kernel area** between the host (Ring -1) and the guest (Ring 0~3)
  - Consumes fewer resources than existing mechanisms
  - Can be loaded at runtime

# Architecture of Shadow-Watcher

- Shadow-watcher
  - Monitors the guest with Light-box
  - Checks if applications of the guest **modify static kernel objects or not** by an event-driven way
    - Code, system table, IDT table, etc
  - Checks the integrity of the guest by **introspecting dynamic kernel objects** by a periodic way
    - **Process list**, LKM list, function pointers of file system and socket

# Detailed Information of Shadow-Box

- If you want to know more about Shadow-box, check the materials below



**black hat**  
ASIA 2017  
MARCH 28-31, 2017  
MARINA BAY SANDS / SINGAPORE

**Myth and Truth about  
Hypervisor-Based Kernel Protector:**  
The Reason Why You Need Shadow-Box

Seunghun Han, Jungwhan Kang  
(hanseunghun || ultract)@nsr.re.kr



**black hat**  
ASIA 2018  
MARCH 20-23, 2018  
MARINA BAY SANDS / SINGAPORE

**Shadow-Box v2:**  
The Practical and Omnipotent Sandbox for ARM

Seunghun Han, Jun-Hyeok Park  
(hanseunghun || parkparkqw)@nsr.re.kr

Wook Shin, Junghwan Kang, HyoungChun Kim  
(wshin || ultract || khche)@nsr.re.kr

#BHASIA / @BlackHatEvents



# Design and Implementation of Gatekeeper Prototype

- Goal is to replace Shadow-watcher functions to Gatekeeper functions!
- Shadow-watcher uses four hardware breakpoints to monitor processes and LKMs
- For Gatekeeper, process and system call information is needed!
  - So, two breakpoints for LKMs are used to monitor system calls

# Key Functions of Linux kernel

- `wake_up_new_task()`
  - Wakes up a new task
- `proc_flush_task()`
  - Removes task from `/proc`
- `entry_SYSCALL_64`
  - Is an entry point of system calls
- `commit_creds()`:
  - Installs new credentials

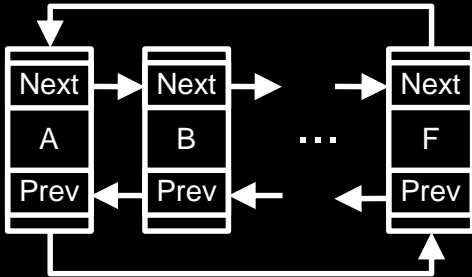
```
ffffffff81aab670 T commit_creds
ffffffff81ab5a20 T wake_up_new_task
ffffffff81cbf9e0 T proc_flush_task
ffffffff822d3e50 T entry_SYSCALL_64
```

# Cred-Related System Calls

```
/* Check cred-related system calls. */  
(syscall_number == __NR_capset)      || (syscall_number == __NR_execve)      ||  
(syscall_number == __NR_keyctl)     || (syscall_number == __NR_prctl)      ||  
(syscall_number == __NR_setfsuid)   || (syscall_number == __NR_setfsuid)   ||  
(syscall_number == __NR_setgid)     || (syscall_number == __NR_setuid)     ||  
(syscall_number == __NR_setregid)   || (syscall_number == __NR_setreuid)   ||  
(syscall_number == __NR_setresgid)  || (syscall_number == __NR_setresuid)  ||  
(syscall_number == __NR_setgroups)  || (syscall_number == __NR_unshare))
```

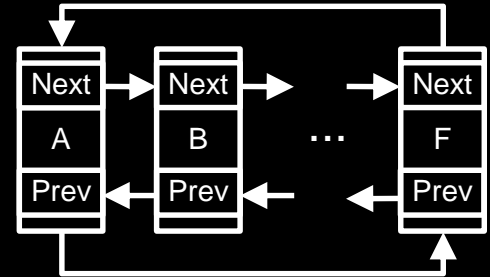
- Only these system calls execute `commit_creds()`
  - Credential information is in the `cred` field of task structure
  - If the applications do not call them, cred **should be unchanged!**

## Task List in the Guest



① Creating initial list with cred info.

## Task List in Gatekeeper



④ Shadowing the list with cred

### Creation Function

```
do_fork()
{
    create_task();
    wake_up_new_task();
}
```

② Inserting H/W breakpoints

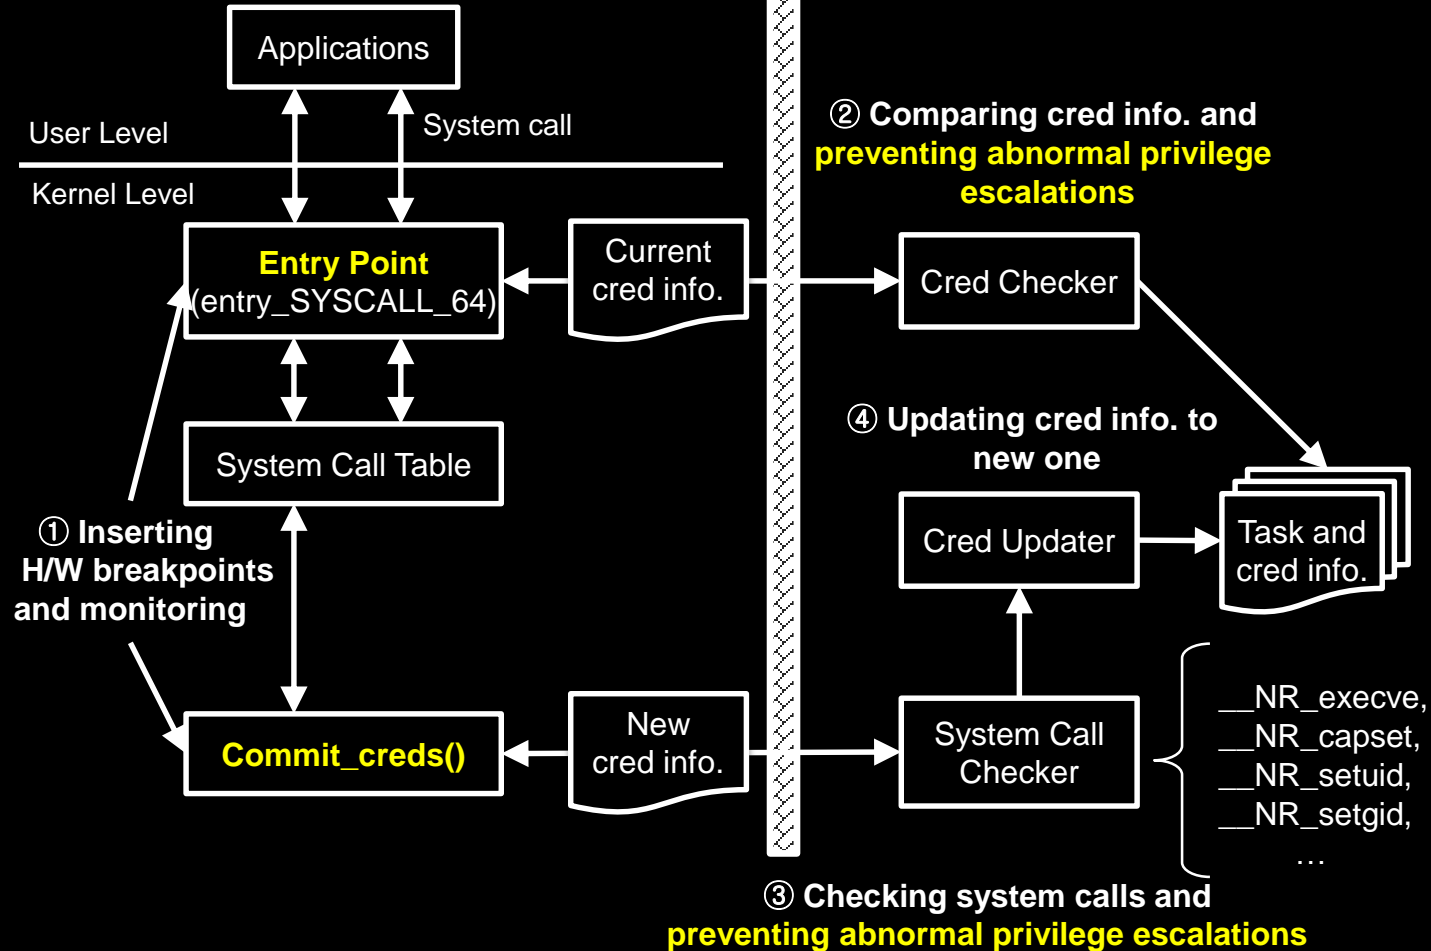
### Deletion Function

```
release_task()
{
    delete_task();
    proc_flush_task();
}
```

③ Monitoring and tracing privilege escalations for Case (3)

## System Calls in the Guest

## Gatekeeper





**Background**

**Analysis of Privilege Escalation  
Exploits**

**Design and Implementation of  
Gatekeeper**  
(with Shadow-Box)

**Evaluation, Demo., Hunting the Beasts**  
(and Why I am Struggling with the Real World Beasts Alone)

**Limitations and Conclusion**

# Evaluation

- Ubuntu 16.04.01 with kernel 4.4.0-21-generic, 4.8.0-41-generic, 4.8.0-58-generic, and 4.10.0-28-generic

NO.	Privilege Escalation CVEs	Detection or Prevention?
1	CVE-2016-4557	Detection
2	CVE-2016-5195	Detection
3	CVE-2016-8655	Prevention
4	CVE-2017-6074	Prevention
5	CVE-2017-7308	Prevention
6	CVE-2017-1000112	Prevention
7	CVE-2017-16995	Prevention



<https://github.com/kkamagui/linux-kernel-exploits>

# GATEKEEPER

Powered by

# SHADOW-BOX

Lightweight Hypervisor-Based Kernel Protector v2.0.0

# DEMO

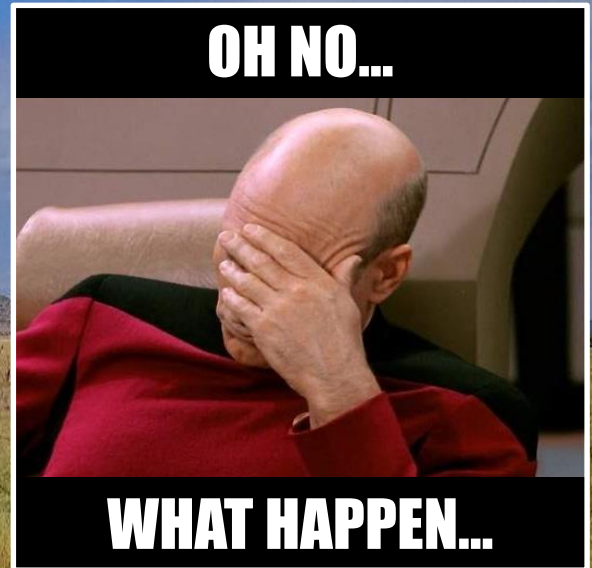
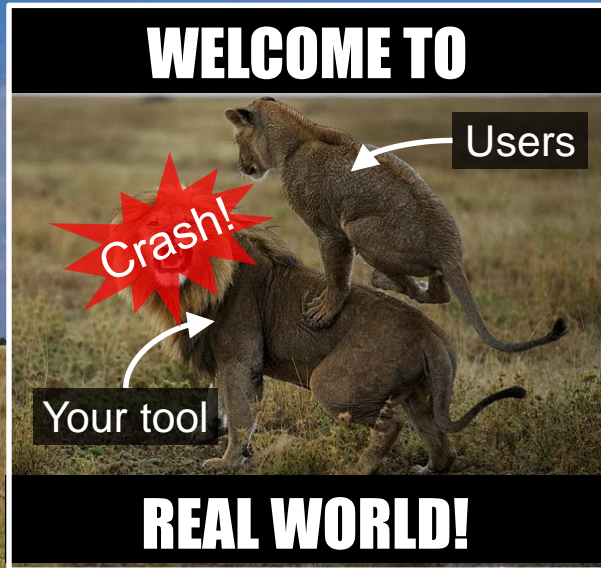
[ 63.077716]  
[ 63.077717]  
[ 63.077718]  
[ 63.077718]  
[ 63.077719]  
[ 63.077719]  
[ 63.077720]  
[ 63.077720]  
[ 63.077721]  
[ 63.077721]  
[ 63.077722]  
[ 63.077723]  
[ 63.077723]  
[ 63.077724]  
[ 63.077725]  
[ 63.077725]  
[ 63.077726]  
[ 63.077726]  
[ 63.077727]  
[ 63.088633]  
[ 63.088634]  
[ 63.102729]  
[ 63.131388]  
[ 63.139345]  
[ 63.284590]  
[ 63.305411]  
[ 63.316548]  
[ 63.337863]  
[ 63.338027]  
[ 63.338028]  
[ 63.338029]  
[ 63.348943]  
[ 63.358808]  
[ 63.359222]  
[ 63.398196]  
[ 63.398197]

shadow-box: CPU Count 4  
shadow-box: Booting CPU 0  
shadow-box: Protect Kernel Code Area  
shadow-box: [\*] Complete  
shadow-box: Framework Preinitialize  
shadow-box: [\*] Complete  
shadow-box: Protect Module Code Area  
shadow-box: [\*] Complete  
shadow-box: Framework Initialize  
shadow-box: [\*] Task count 425  
shadow-box: [\*] Module count 118  
shadow-box: [\*] Complete  
shadow-box: Lock IOMMU  
shadow-box: [\*] Intel Integrated Graphics is detected  
shadow-box: [\*] Lock IOMMU complete  
shadow-box: Execution Complete  
shadow-box: **errorcode=0**



# Deployment, Real World, and the Beasts

Real world is Serengeti and...  
**BEASTS live there!!!**



# There were many friends...

## SecVisor: A Tiny Hypervisor to Provide Lifetime Kernel Code Integrity for Commodity OSes

Arvind Seshadri  
Cylab/CMU  
Pittsburgh, PA, USA  
arvind@cs.cmu.edu

Mark Luk  
Cylab/CMU  
Pittsburgh, PA, USA  
mluk@ece.cmu.edu

Ning Qu  
Cylab/CMU  
Pittsburgh, PA, USA  
quning@cmu.edu

Adrian Perrig  
Cylab/CMU  
Pittsburgh, PA, USA  
perrig@cmu.edu

### ABSTRACT

We propose SecVisor, a tiny hypervisor that ensures code integrity for commodity OS kernels. In particular, SecVisor ensures that only user-approved code can execute in kernel mode over the entire system lifetime. This protects the kernel against code injection attacks, such as kernel rootkits. SecVisor can achieve this property even against an attacker who controls everything but the CPU, the memory controller, and system memory chips. Further, SecVisor can even defend against attackers with knowledge to directly tamper with kernel registers. Our goal is to make SecVisor amenable to commodity OSes.

# 2007

### 1. INTRODUCTION

Computing platforms are steadily increasing in complexity, incorporating an ever-growing range of hardware and supporting an ever-growing range of applications. Consequently, the complexity of OS kernels is steadily increasing. The increased complexity of OS kernels also increases the risk.

The effect of these vulnerabilities is that, despite many efforts to make hardware more secure, today's systems are increasingly vulnerable to attacks that exploit these vulnerabilities. This kernel mode rootkit is a privilege escalation attack.



## NumChecker:

## A System Approach for Kernel Rootkit Detection and Identification

Xueyang Wang, Ph.D.

Xiaofei (Rex) Guo, Ph.D.

(xueyang.wang, guo@cs.cmu.edu) @PAM\* intel.com

# 2016

## Lares: An Architecture for Secure Active Monitoring Using Virtualization

Bryan D. Payne Martim Carbone Monirul Sharif Wenke Le  
School of Computer Science  
Georgia Institute of Technology  
Atlanta, Georgia 30332-0765  
{bdpayne, mcarbone, msharif, wenke}@cc.gatech.edu

### Abstract

Host-based security tools such as anti-virus and intrusion detection systems are not adequately protected on today's computers. Malware is often designed to immediately disable any security tools upon installation, rendering them useless. While current research has focused on moving these vulnerable security tools into virtual machines, this approach cripples the security tool.

ing a full-featured anti-virus, intrusion detection, or intrusion prevention system. Previous efforts to implement these types of systems within a protected VM have resorted to implementing the systems with crippled functionality. What was missing in these systems was the ability to do active monitoring. Active monitoring is when the security tool places a hook inside the system being monitored. When the system is inside the hook, it will interrupt execution and send data to the security tool. Active monitoring can also

# 2008

## Guest-Transparent Prevention of Kernel Rootkits with VMM-based Memory Shadowing

Ryan Riley  
Purdue University  
rileyrd@cs.purdue.edu

Xuxian Jiang  
George Mason University  
xjiang@gmu.edu

Dongyan Xu  
Purdue University  
dxu@cs.purdue.edu

### Abstract

Kernel rootkits pose a significant threat to computer systems as they run at the highest privilege level and have unrestricted access to the resources of their victims. Many current efforts in kernel rootkit defense focus on the detection of kernel rootkits – after a rootkit attack has taken place, while the smaller efforts in kernel rootkit prevention exhibit limitations in their capability or deployability. In this paper, we present a kernel rootkit prevention system called NICKLE which addresses a common, but often overlooked, characteristic of most kernel rootkits: the need for executing their own kernel code. NICKLE, a lightweight, virtual machine monitor (VMM)-based system, consequently prevents unauthorized kernel execution for unmodified OSes. NICKLE is based on a new scheme

# 2008

## Ensuring Operating System Kernel Integrity with OSck

Ogden S. Hoffmann Alan M. Dunn Sangman Kim Indrajit Roy\* Emmett Witchel  
The University of Texas at Austin \*HIP Labs  
{os\_h, adunn, sangman, witchel}@cs.utexas.edu indrajit@hip.com

### Abstract

Kernel rootkits that modify operating system state to avoid detection are a dangerous threat to system security. This paper presents OSck, a system that discovers kernel rootkits by detecting malicious modifications to operating system data. OSck integrates and extends existing techniques for detecting rootkits, and verifies safety properties for large portions of the kernel heap with minimal overhead. We deduce type information for verification by analyzing unmodified kernel source code and in-memory kernel data.

High performance integrity checks that execute concurrently with a running operating system must detect, and be able to respond to, a deterministic solution for ensuring kernel integrity. We introduce OSck, a security tool that ensures kernel integrity by verifying the state of operating system data structures in order to guarantee the integrity of the operating system.

change the state of operating system data structures in order to guarantee the integrity of the operating system. OSck detects when the state of kernel data structures violates the integrity properties specified to the hypervisor. If a rootkit compromises kernel integrity, the hypervisor can take appropriate action, such as terminating the operating system or alerting an administrator.

We extend previous work in hypervisor-based monitoring in four important directions: 1. OSck verifies type safety properties for the kernel heap through a series of memory safety tests (covering a heap overflow graph). This approach is based on extracting annotations about memory usage from the kernel source code. 2. OSck verifies kernel integrity by verifying the state of operating system data structures in order to guarantee the integrity of the operating system. 3. OSck verifies kernel integrity by verifying the state of operating system data structures in order to guarantee the integrity of the operating system.

# 2011

# And... Gone with the Laboratory!

## SecVisor: A Tiny Hypervisor to Provide Lifetime Kernel Code Integrity for Commodity OSes

Arvind Seshadri  
Cylab/CMU  
Pittsburgh, PA, USA  
arvind@cs.cmu.edu

Mark Luk  
Cylab/CMU  
Pittsburgh, PA, USA  
mluk@ece.cmu.edu

Ning Qu  
Cylab/CMU  
Pittsburgh, PA, USA  
quining@cmu.edu

### ABSTRACT

We propose SecVisor, a tiny hypervisor that ensures code integrity for commodity OS kernels. In particular, SecVisor ensures that only user-approved code can execute in kernel mode over the entire system lifetime. This protects the kernel against code injection attacks, such as kernel rootkits. SecVisor can achieve this property even against unpatchable commodity OSes, but the OS must be manually customized and system updates require that the OS be re-approved by SecVisor. SecVisor can be deployed on commodity OSes and can even defend against attacks with kernel rootkits that are not patched.

2007

### 1. INTRODUCTION

Computing platforms are an ever-growing range of applications of OS kernels is steadily increasing. The effect of these rootkits is to increase the risk of system compromise. SecVisor can be deployed on commodity OSes and can even defend against attacks with kernel rootkits that are not patched.

## Lares: An Architecture for Secure Active Monitoring Using

Bryan D. Payne Martin Carbone Monirul Sharif Wen  
School of Computer Science  
Georgia Institute of Technology  
Atlanta, Georgia 30332-0765  
{bdpayne, mcarbone, msharif, wenke}@cc.gatech.edu

### Abstract

Host-based security tools such as anti-virus and intrusion detection systems are not adequately protected on today's computers. Malware is often designed to im-

ing a full-fledged anti-virus or intrusion detection system prevention system. Previous efforts to implement these types of systems within a protected VM have resorted to implementing the systems with crippled functionality. What was missing in these systems was the ability to do active

monitoring. Active monitoring is the ability to detect and respond to attacks before the system being monitored. When a host-based security tool is compromised, it will no longer be able to detect and respond to attacks. Active monitoring can

2008

## Guest-Transparent Prevention of Kernel Rootkits with VMM-based Memory Shadowing

Xijiang Jiang  
George Mason University  
xjiang@gmu.edu

Dongyan Xu  
Purdue University  
dxu@cs.purdue.edu

### Abstract

A significant threat to computer systems as they run at the highest privilege level is the resources of their victims. Many current efforts in kernel rootkit detection of kernel rootkits – after a rootkit attack has taken place, while the smaller kernel rootkit prevention exhibit limitations in their capability or deployability. In this paper, we propose a new system called NIKKAR which addresses a common, but often overlooked kernel rootkit: the rootkit for modifying their own kernel code. NIKKAR is a VMM-based system that can transparently prevent unauthorized kernel modifications. NIKKAR is based on a new scheme

2008

## Operating System Kernel Integrity with OSck

Alan M. Dunn Sangnam Kim Indrajit Roy\* Emmett Witchel  
The University of Texas at Austin \*HIP Labs  
{am.dunn, sangnam.kim, indrajit.roy, ewitchel}@cs.utexas.edu

Operating system state to avoid detection of unauthorized access to computer resources and to prevent detection of kernel rootkits. OSck integrates and extends existing techniques for detecting rootkits, and verifies safety properties for large portions of the kernel heap with minimal overhead. We deduce type information for verification by analyzing unmodified kernel source code and in-memory kernel data

change the state of operating system data structures in order to gain unauthorized access to computer resources and to prevent detection. OSck detects when the state of kernel data structures violates the integrity properties specified to the hypervisor. If a rootkit compromises kernel integrity, the hypervisor can take appropriate action, such as terminating the operating system or alerting an administrator.

We extend previous work in hypervisor-based monitoring in four important directions:

1. OSck enables dynamic verification of the kernel heap through a VMM-based system, rather than requiring a hypervisor-based patch. This approach is based on monitoring operations that change the state of kernel data structures. OSck can detect when the state of kernel data structures violates the integrity properties specified to the hypervisor. If a rootkit compromises kernel integrity, the hypervisor can take appropriate action, such as terminating the operating system or alerting an administrator.

2011

NOOOO, BUDDY!



DON'T LEAVE ME ALONE!



# Beasts and Hunting – Easy Level



- **(1) Code area was not immutable!**
  - I turned CONFIG\_JUMP\_LABEL off in kernel config
- **(2) Cache types of memory areas were misconfigured!**
  - I set uncacheable type by default and write-back type to “System RAM” area according to /proc/iomem

# Beasts and Hunting – Normal Level



- (1) Machine check error (MCE) exception could not handled in the host (Ring -1)!
  - I enabled MCE polling feature in the kernel!
  - **I found a kernel bug (!?)** and contributed a patch to kernel
- (2) New instructions such as XSAVES and XRSTORS were added to new CPUs!
  - I added new instruction handlers according to Intel manuals

# Beasts and Hunting – Hardcore Level

- **MELTDOWN** was revealed!

- And Page Table Isolation (PTI) feature was added to the kernel!
  - PTI flushes kernel-level page tables every privilege-level transition! (ring 0  $\leftarrow$   $\rightarrow$  ring 3)
- I read the meltdown paper and kernel source code **AGAIN** and **AGAIN**!
- Finally, I could change memory mapping of Shadow-box according to PTI patches

**I AM ALWAYS ANGRY!**



**NUUUUTS!!!**

<https://github.com/kkamagui/shadow-box-for-x86/issues>





**Background**

**Analysis of Privilege Escalation  
Exploits**

**Design and Implementation of  
Gatekeeper**  
(with Shadow-Box)

**Evaluation, Demo., Hunting the Beasts**  
(and Why I am Struggling with the Real World Beasts Alone)

**Limitations and Conclusion**

# Limitations of Gatekeeper Prototype

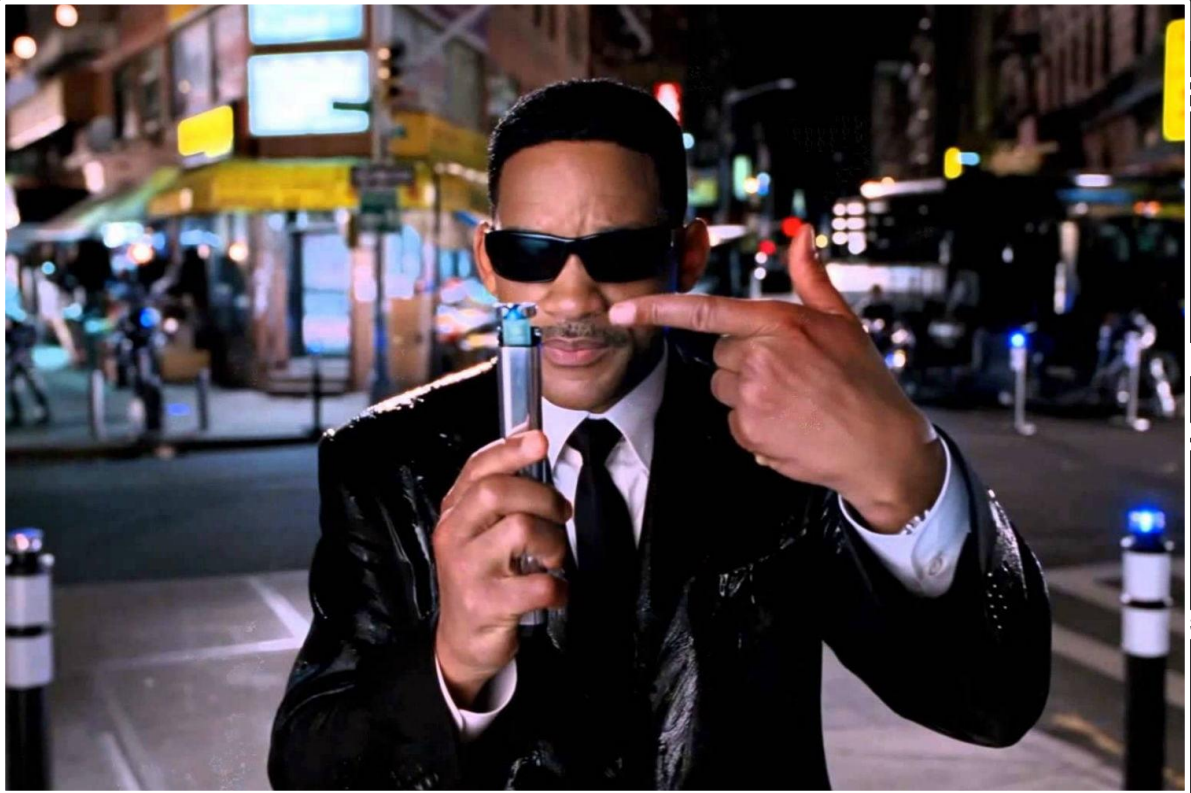
- Gatekeeper only works in 64bit mode
  - It cannot monitor system calls with 32bit mode (compat) and interrupt mode
- Gatekeeper's prevention feature can be avoided
  - If you make complicate ROP chain to change privilege objects (not cred) in user level
- Gatekeeper can be neutralized
  - If cred-related system calls have vulnerabilities



# Conclusion

- Three types of privilege escalation exploits can be detected and prevented by **Gatekeeper (based on Shadow-box)**
- **Shadow-box is the last man standing** and struggling with the real world beasts alone (unfortunately)
- **Exploits might evolve** to change privileged objects, not cred (because of this presentation)

# Conclusion



# Questions?



## CONTRIBUTION!

Project : <https://github.com/kkamagui/shadow-box-for-x86>

Contact: hanseunghun@nsr.re.kr, @kkamagui1

# Reference

- “Myth and Truth about Hypervisor-Based Kernel Protector: The Reason Why You Need Shadow-Box”, Black Hat Asia 2017
- “Shadow-Box: The Practical and Omnipotent Sandbox”, HITBSecConf 2017
- “Shadow-Box v2: The Practical and Omnipotent Sandbox for ARM”, Black Hat Asia 2018
- “Proposal of a Method to Prevent Privilege Escalation Attacks for Linux Kernel”, Linux Security Summit 2017
- Shadow-box and Gatekeeper, <https://github.com/kkamagui/shadow-box-for-x86>
- Linux Kernel Exploits, <https://github.com/kkamagui/linux-kernel-exploits>